



Cassandra Performance Tuning

Apache Cassandra:
Environment Tuning

Learning objectives

- **Discuss environment tuning outside of Cassandra**
- Identify tools and tuning strategies for the JVM
- Discuss working with page cache
- Identify tools to monitor the CPU

What are the most common bottlenecks in Cassandra?

Common Bottlenecks

- Inadequate hardware
- Insufficient or incorrect memory cache tuning
- Poorly configured JVM parameters
- High CPU utilization

Question

What is the best way to cope with inadequate node hardware in a Cassandra cluster?

How do you cope with inadequate node hardware?

- Upgrade the hardware
- Analyze your operating system

What are some of the technologies upon which a Cassandra node depends?

Java technologies and tools

- Java—an object-oriented, current and class-based programming language.
- JVM—a process virtual machine that can operate Java bytecode.
- JNA—(Java Native Access) an API for allowing Java programs access to shared libraries.
- Heap—an area of memory used for dynamic memory allocation.
- JMX—(Java Management Extensions) supplies tools for managing and monitoring resources.

Learning objectives

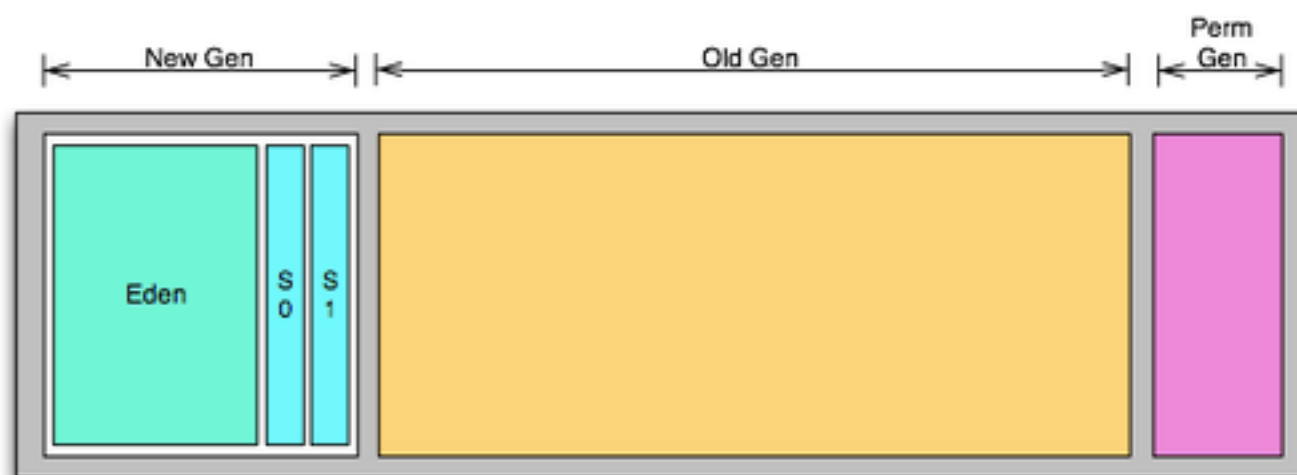
- Discuss environment tuning outside of Cassandra
- **Identify tools and tuning strategies for the JVM**
- Discuss working with page cache
- Identify tools to monitor the CPU

JVM and Garbage Collection (GC)

- Recall that garbage collection is a process that a JVM is undergoing all the time to clean out any processes that are not live
 - Objects are moved and deleted to free up memory
 - GC should happen often enough to create larger “holes” of free memory, but not so often that the CPU is churning on GC all the time
- Since Cassandra runs in a JVM, the pauses to do garbage collection affect Cassandra’s performance
 - Sizing the JVM is important to performance
 - The number of CPUs can also affect performance

JVM available memory

- New generation—broken up into eden and survivor spaces
 - When eden fills with new object, minor garbage collection occurs
 - Application pauses when ParNew collector is run
- Old generation
 - Contains objects that have survived long enough to not be collected by gc
 - When a certain percentage (75% default) is full, CMS collector is run



Graphic by Blake Eggleston

JVM heap options and Cassandra performance

- Setting *MAX_HEAP_SIZE* to not more than 8GB
 - Large heaps can introduce GC pauses that lead to latency
 - On the other hand, different workloads can justify different settings
 - SHIFT, for instance, found that using 10GB improved their performance
- Setting to *HEAP_NEWSIZE* to 100MB per core
 - 8 cores would mean 800MB
 - But again, different workloads may react differently
- Interesting to set these values to low settings and high settings and compare the results
 - In the lab for this section, you will do just that, so you can see for yourself.

Using *nodetool tpstats* to discover JVM issues

- Displays the number of active, pending, and completed tasks for each of the thread pool that Cassandra uses for stages of operations


```
ubuntu@ip-172-31-8-219:~$ nodetool -h node0 tpstats
```

Pool Name	Active	Pending	Completed	Blocked	All time blocked
ReadStage	0	0	4681361	0	0
RequestResponseStage	0	2	11269852	0	0
MutationStage	0	0	18235083	0	0
ReadRepairStage	0	0	111456	0	0
ReplicateOnWriteStage	0	0	0	0	0
GossipStage	0	0	4775637	0	0
AntiEntropyStage	0	0	187096	0	0
MigrationStage	0	0	70	0	0
MemoryMeter	0	0	342	0	0
MemtablePostFlusher	0	0	82753	0	0
FlushWriter	0	0	17698	0	1
MiscStage	0	0	53456	0	0
PendingRangeCalculator	0	0	3	0	0
commitlog_archiver	0	0	0	0	0
AntiEntropySessions	0	0	3341	0	0
InternalResponseStage	0	0	53459	0	0
HintedHandoff	0	0	20	0	0

Message type	Dropped
RANGE_SLICE	0
READ_REPAIR	0
PAGED_RANGE	0
BINARY	0
READ	0
MUTATION	0
_TRACE	2690912
REQUEST_RESPONSE	0
COUNTER_MUTATION	0

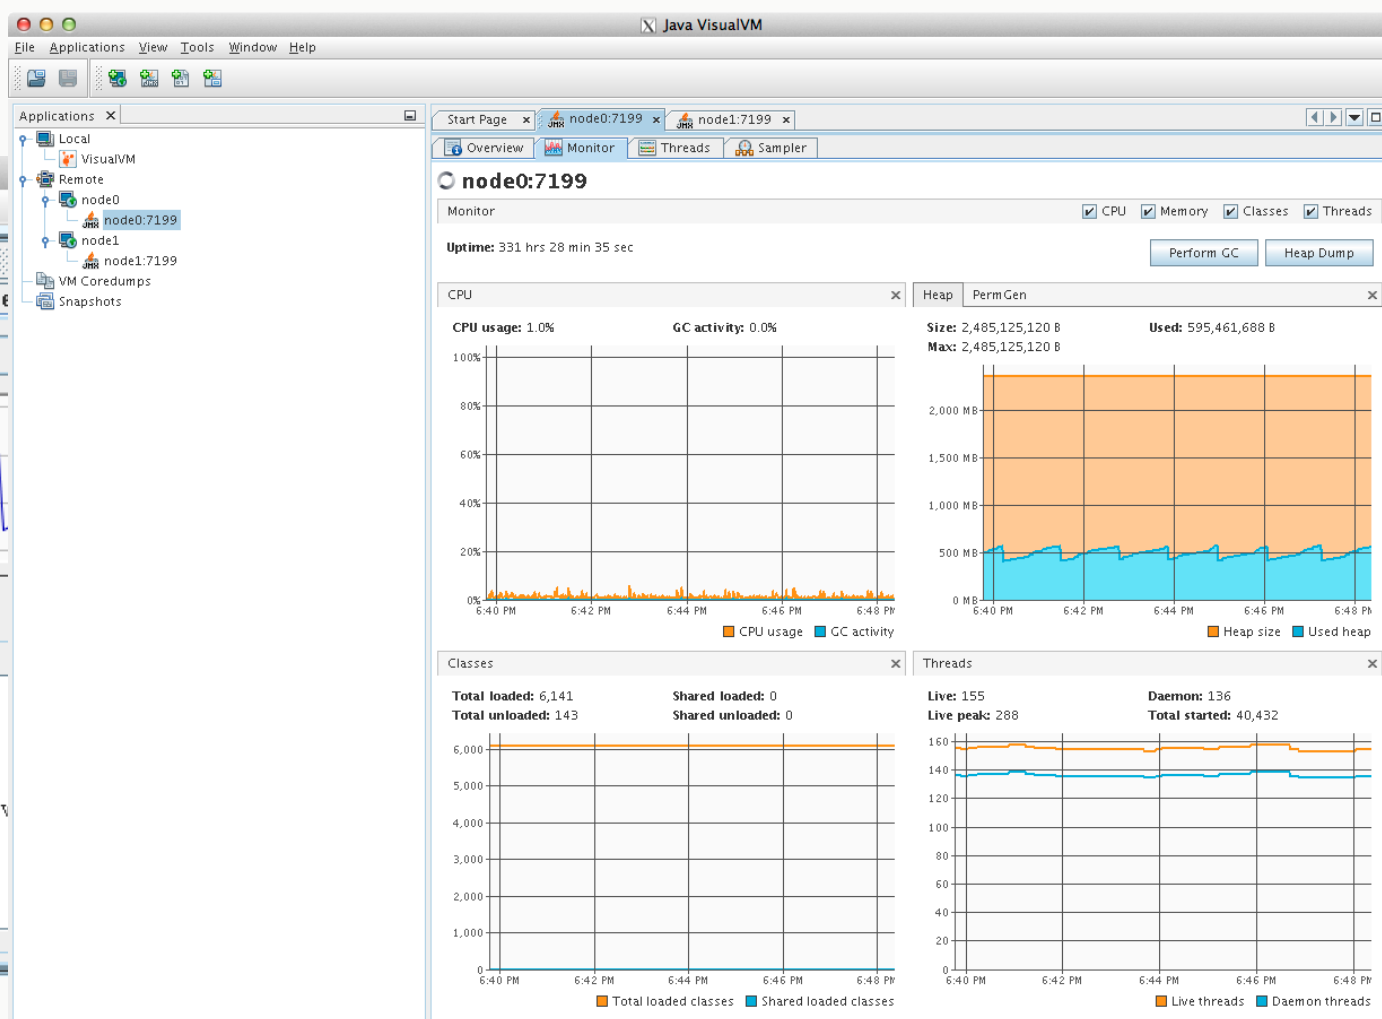
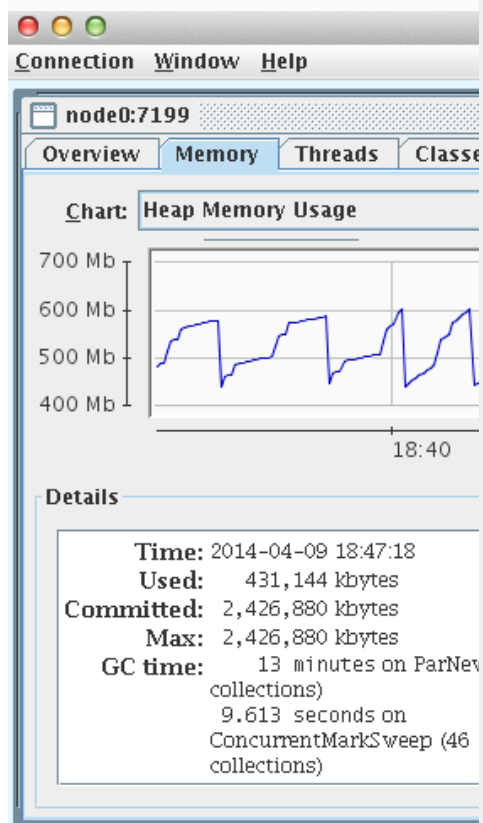
Using OpsCenter to explore JVM issues

- *OpsCenter* can provide visual insight into the JVM performance with the following graphs:
 - JVM Collection Count for both ParNew and CMS
 - JVM Collection Time for both ParNew and CMS
 - Heap Max, Heap Used, Heap Committed
 - Number of pending tasks – drill down for more information



Using JMX clients to explore JVM issues

- A couple of visual tools that can be used to “watch” heap action
 - jconsole
 - jvisualvm



Demo 1: Viewing JVM heap memory using jconsole and OpsCenter

Exercise I: Tuning JVM heap

Learning objectives

- Discuss environment tuning outside of Cassandra
- Identify tools and tuning strategies for the JVM
- **Discuss working with page cache**
- Identify tools to monitor the CPU

What is page caching?

- Accelerates access to files on non-volatile storage
 - When reading and writing to hard disk, Linux also stores data in unused areas of memory.
 - These act as a cache, and allow that data to be quickly read from memory if they are accessed again.
- Page cache is sometimes called Buffer cache
 - At one time, page cache and buffer cache were separate, but now they have been combined.
- The command `free -m` can be used to discover how much RAM is used for page cache

```
ubuntu@ip-172-31-8-219:~$ free -m
```

	total	used	free	shared	buffers	cached
Mem:	7450	7125	324	0	125	6583
-/+ buffers/cache:		417	7033			
Swap:	0	0	0			

How does Cassandra utilize page caching?

- Linux will cache most recently used disk pages.
- The more RAM you have, the more data that will be held in memory.
 - Once memory is filled up, the most stale data will be overwritten.
- **mmap makes file access very efficient for C***
 - Using file mapping, the kernel maps your program's virtual pages directly onto the page cache.
- **Page cache makes writes more efficient**
 - OS flushes page cache to disk.
- **Page cache makes reads more efficient.**
 - OS can access reads from memory more quickly than from disk.
- **How can page caching improve performance?**
 - More RAM, more page cache that is available.

How do you triage root cause for out of memory (OOM) errors?

- **Writing to partitions with inadequate timeout**
 - Writes back up in memory if the throughput is too slow.
 - Increase the disk throughput with local disk and/or SSDs.
- **Reading from partitions with lots of deletes, including TTLs**
 - Reads require ALL the data, including the deleted data, to be read into memory to get the queried data.
 - If you must use frequent TTLs and deletes, consider whether or not you can cluster these to fewer rows – deleting a row can solve many of the issues associated with this.
- **Client-side joins**
 - Too much data must be resident on memory to accomplish the joins.
 - Try to eliminate client-side joins as much as possible – consider redesigning your schemas.

Exercise 2: Working with the Linux page cache

Learning objectives

- Discuss environment tuning outside of Cassandra
- Identify tools and tuning strategies for the JVM
- Discuss working with page cache
- **Identify tools to monitor the CPU**

What operations benefit from better CPUs?

- Garbage collection is CPU-intensive
 - Add more CPUs or faster CPUs, and garbage collection will run faster.
- Compression is also CPU-intensive
 - Just like GC, compression stresses the CPUs.

How do you triage root cause for CPU errors?

- Check CPU usage with *dstat*
- Check CPU usage with *OpsCenter*
 - CPU graphs can be added to monitor CPU usage
- What to do?
 - Add nodes
 - Use nodes that have more and faster CPUs

Demo 2: Using command line tools to monitor CPU

Review Questions

- How can insufficient hardware affect performance?
- What impact does garbage collection have on performance?
- Where can the page cache be configured?
- What Cassandra operations are CPU intensive?
- What tools are available for monitoring the CPU?

